



CSC412

Assignment 3

Vibhavi Peiris

Student #: 1000597687

April 6, 2018

1. Problem 1: L2-Regularized Logistic Regression

a. Use code from A2 with 300 training points

```
#A3 Q1A From my-A2 Q3C code
def one_per_class(images, labels):
    out_images = np.zeros((10,images.shape[1]))
    out_labels = np.zeros((10,10))
    classes = np.where(labels == 1)[1] # get the class digit for each image by getting column idx of ones in labels
    #get first image in training set with each class label
    for i in range(0,10):
        img_num = np.where(classes == i)[0][0]
        out_images[i,:] = images[img_num,:]
        out_labels[i,:] = labels[img_num,:]
    return out_images,out_labels

def cost_function(w):
    sum_final = 0 #temporary create sum_final var
    dem = logsumexp(np.dot(np.transpose(w),grad_images))

    #muticlass likelihood function is sum from 0 to k of label*predictive_log_likelihood
    for k in range(0,10):
        log_pc_x = np.dot(np.transpose(w[:,k]),grad_images) - dem
        if k == 0:
            sum_final = np.dot(grad_labels[k],log_pc_x)
        else:
            sum_final = sum_final + np.dot(grad_labels[k],log_pc_x)
    return sum_final

def logistic_gradient_desc(iterations,lr):
    #set globals so that cost function can access these values after usign autograd w.r.t. w
    global current_c
    global grad_images
    global grad_labels

    w = np.zeros((784,10)) #create the weights
    for i in range(0,iterations):
        for img_num in range(0,new_images.shape[0]):
            #get gradient of cost function/likelihood
            grad_images = new_images[img_num,:] #get current image
            grad_labels = new_labels[img_num,:] #get labels for current image
            current_c = img_num #sinces we sampled 1 image for each class in order c = img_num
            cost_grad = elementwise_grad(cost_function)

            #update weights
            w = w + lr*cost_grad(w)
        print(i)
    return w

#run A2 Q3C code
new_images = train_images
new_labels = train_labels
grad_images=grad_labels = new_images #temporary just to create a gobal var for use with autograd
current_c = 0 #temporary just to create a gobal var for use with autograd
weights = logistic_gradient_desc(1000, 0.01) #5000 iterations with a common learning rate of 0.01
save_images(np.transpose(weights),'Q1a')
```

```

#A2 Q3d code
def avg_pred_log(w,images):
    log_pc_x = 0
    for i in range(0,images.shape[0]):
        current_log_pc_x = np.dot(np.transpose(w),images[i,:]) - logsumexp(np.dot(np.transpose(w),images[i,:]))
        log_pc_x = log_pc_x + current_log_pc_x

    return np.sum(log_pc_x)/float(images.shape[0])

def predict_regression(images, w):
    predictions = np.zeros((images.shape[0],w.shape[1])) #N by 10

    #find best class true class for each image
    for i in range(0,images.shape[0]):
        best_class = np.argmax(np.dot(np.transpose(w),images[i,:])) #choose the class with highest
        predictions[i,best_class] = 1 #set index = digit to 1 as it is the best prediction for current image

    return predictions
def Q3D_report(train_images,train_labels,test_images,test_labels,w):
    #average for train
    avg_likelihood_train = avg_pred_log(w,train_images)
    print("Avg train log likelihood ",avg_likelihood_train)


    #average for test
    avg_likelihood_test = avg_pred_log(w,test_images)
    print("Avg test log likelihood ",avg_likelihood_test)

    #predictions for train
    predict_train = predict_regression(train_images, w)
    total_correct_train = np.sum(np.nonzero(predict_train)[1] == np.nonzero(train_labels)[1]) #get total number of correct predictions
    accuracy_train = total_correct_train/float(train_labels.shape[0]) #get accuracy
    print('Train Accuracy ',accuracy_train)

    #predictions for test
    predict_test = predict_regression(test_images, w)
    total_correct_test = np.sum(np.nonzero(predict_test)[1] == np.nonzero(test_labels)[1]) #get total number of correct predictions
    accuracy_test = total_correct_test/float(test_labels.shape[0]) #get accuracy
    print('Test Accuracy ',accuracy_test)

#run A2 Q3D code
Q3D_report(train_images,train_labels,test_images,test_labels,weights) #REPORT FOR A3 Q1a

```



Average train log likelihood -117.989022988

Average test log likelihood -97.3003729208

Train Accuracy 1.0

Test Accuracy 0.7727

b. MAP Estimator

$$N(w_{cd} | 0, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-w^2}{2\sigma^2}\right)$$

$$\log(N(w_{cd} | 0, \sigma^2)) = \log\left(\frac{1}{\sqrt{2\pi\sigma^2}}\right) - \frac{w^2}{2\sigma^2} = -\log(\sqrt{2\pi\sigma^2}) - \frac{w^2}{2\sigma^2}$$

$$\begin{aligned} \log(p(t | X, w) p(w | \sigma^2)) &= \log\left(\prod_{i=0}^{300} \left(\frac{\exp(w_t^T x_i)}{\sum_{c=0}^9 \exp(w_c^T x_i)}\right) \prod_{c=0}^9 \prod_{d=0}^{784} N(w_{cd} | 0, \sigma^2)\right) \\ &= \sum_{i=0}^{300} \log\left(\frac{\exp(w_t^T x_i)}{\sum_{c=0}^9 \exp(w_c^T x_i)}\right) + \sum_{c=0}^9 \sum_{d=0}^{784} \log\left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(\frac{-w^2}{2\sigma^2}\right)\right) \\ &= \sum_{i=0}^{300} \left(w_t^T x_i - \log \sum_{c=0}^9 \exp(w_c^T x_i)\right) + \sum_{c=0}^9 \sum_{d=0}^{784} -\log(\sqrt{2\pi\sigma^2}) - \frac{w^2}{2\sigma^2} \end{aligned}$$

$$\begin{aligned} \nabla_w \log(p(t | X, w) p(w | \sigma^2)) &= \nabla_w \sum_{i=0}^{300} \left(w_t^T x_i - \log \sum_{c=0}^9 \exp(w_c^T x_i)\right) + \nabla_w \sum_{c=0}^9 \sum_{d=0}^{784} -\log(\sqrt{2\pi\sigma^2}) - \frac{w^2}{2\sigma^2} \\ &= \sum_{i=0}^{300} \left(x_i - \frac{\exp(w_c^T x_i) x_i}{\sum_{c=0}^9 \exp(w_c^T x_i)}\right) - \frac{w}{\sigma^2} \end{aligned}$$

c. Fit map

```
#A3 Q1C

#logistic regression with gradient descent using map
def grad_desc(iterations,lr,sigma):
    #set globals so that cost function can access these values after usign autograd w.r.t. w
    global grad_images
    global grad_labels

    w = np.zeros((784,10)) #create the weights
    for i in range(0,iterations):
        for img_num in range(0,new_images.shape[0]):

            #get gradient of cost function/likelihood
            grad_images = new_images[img_num,:] #get current image
            grad_labels = new_labels[img_num,:] #get labels for current image
            cost_grad = elementwise_grad(cost_function)

            #update weights
            w = w + lr*cost_grad(w)

        #NEW ADDITION FOR A3
        w = w - w/sigma**2

    print(i)
    return w

#run cod for A3 q1c
print("Map logistsic regression")

##Testing for best sigma value was 36
#for i in range(1,10):
#    #sigma = i**2 # from 5 to 100
#    #print(sigma)
#    #map_weights = grad_desc(100, 0.01,sigma) #5000 iterations with a common learning rate of 0.01
#    #save_images(np.transpose(map_weights),'Q1c')
#    #Q3D_report(train_images,train_labels,test_images,test_labels,map_weights)

#sigma = 1/i**2 #from 1 to 1/100
#print(sigma)
#map_weights = grad_desc(100, 0.01,sigma) #5000 iterations with a common learning rate of 0.01
#save_images(np.transpose(map_weights),'Q1c')
#Q3D_report(train_images,train_labels,test_images,test_labels,map_weights)
sigma = 36
map_weights = grad_desc(1000, 0.01,sigma) #5000 iterations with a common learning rate of 0.01
save_images(np.transpose(map_weights),'Q1c')
Q3D_report(train_images,train_labels,test_images,test_labels,map_weights)
```


Average train log likelihood -93.0039030291

Average test log likelihood -77.2325259507

Train Accuracy 1.0

Test Accuracy 0.7711

Best sigma value was 36. With 4,9,16,25 close behind.



2. Problem 2: Bayesian Logistic Regression using Stochastic Variational Inference

a. Number of parameters?

For w it is $784 * 10 = 7840$ parameters.

For ϕ its (mean + standard deviation) $= 7840 + 7840 = 14960$ parameters.

b. Code SVI.

```
def elbo_estimate(var_params, logprob, num_samples, rs):
    """Provides a stochastic estimate of the variational lower bound.
    var_params is (mean, log_std) of a Gaussian."""
    mean, log_std = var_params
    samples = sample_diag_gaussian(mean, log_std, num_samples, rs)
    log_ps = logprob(samples)
    log_qs = diag_gaussian_log_density(samples, mean, log_std)
    E_q = np.sum(log_ps - log_qs) / num_samples # E_q(z|x)[log p(x,z) - log q(z|x)]
    return E_q

def logprob_given_data(params):
    data_logprob = logistic_logprob(params, train_images, train_labels)
    prior_logprob = np.sum(np.sum(-np.log(np.sqrt(2*np.pi*prior_std)) - (params**2)/(2*prior_std), axis=2), axis=1)
    return data_logprob + prior_logprob
```

c. Compute accuracy for test

```
predict_test = predict_regression(test_images, np.transpose(optimized_params[0])) #A3 q1/A2 code
total_correct_test = np.sum(np.nonzero(predict_test)[1] == np.nonzero(test_labels)[1]) #get total number of correct predictions
accuracy_test = total_correct_test / float(test_labels.shape[0]) #get accuracy
print("\nTest Accuracy ", accuracy_test)


##testing a bunch of std values, std = 1 is best with 77.68% and std =9 is second best 77.09%
for i in range(1,10):
    #prior_std = i**2 # from 5 to 100
    #print(prior_std)
    #optimized_params = adam(objective_grad, init_params, step_size=0.05, num_iters=100, callback=print_perf)
    #predict_test = predict_regression(test_images, np.transpose(optimized_params[0])) #A3 q1/A2 code
    #total_correct_test = np.sum(np.nonzero(predict_test)[1] == np.nonzero(test_labels)[1]) #get total number of correct predictions
    #accuracy_test = total_correct_test / float(test_labels.shape[0]) #get accuracy
    #print("\nTest Accuracy ", accuracy_test)

    #prior_std = 1/i**2 #from 1 to 1/100
    #print(prior_std)
    #optimized_params = adam(objective_grad, init_params, step_size=0.05, num_iters=100, callback=print_perf)
    #predict_test = predict_regression(test_images, np.transpose(optimized_params[0])) #A3 q1/A2 code
    #total_correct_test = np.sum(np.nonzero(predict_test)[1] == np.nonzero(test_labels)[1]) #get total number of correct predictions
    #accuracy_test = total_correct_test / float(test_labels.shape[0]) #get accuracy
    #print("\nTest Accuracy ", accuracy_test)
```

Standard deviation = 1 is the best with 77.68% and Standard deviation = 9 is second best 77.09%. These number are only slightly higher than MAP inferences accuracy.


d. *Plot 10 images*

i. *Mean*




The mean is as I expected because the darkest regions would be the most common pixels of a digit (middle for 1).

ii. *Variation*



The variation would have the opposite effect of mean because the further away (up to a certain distance) from the digits mean locations the more variation on pixel values. I.E. for digit 1 the close pixels around the one would have high variation (due to all possible ways one is drawn)

iii. *Samples*



The sample is not what I expected, there is more noise. I expected it to be a combination of mean and variance.

e. *Single sample q*

When $x_d \in B$ that means that $x_d = 0$. So $w_{cd}^T x_d$ will be 0. Which means it does not affect the optimal because $p(t|w,x)$ is 0 (substitute $w_{cd}^T x_d = 0$ into the $p(t|w,x)$ equation).


f. *Posterior questions*

i. *Does training affect posterior and variational?*


No, true posterior seems to stay the same through the training. Yes, Variational posterior changes.

ii. *How does standard deviation affect posterior?*


Increasing standard deviation makes the true posterior wider across the horizontal.



Standard deviation = 1



standard deviation = 10




standard deviation = 100

iii. *Px2's affect?*

It makes the true posterior oval shaped & different from variational posterior.

iv. *Changing px2's affect?*

It makes true posterior more circular similar to variational.



v. *Using improper flat model?*

No